

# — Under the Hood of the Google App Engine —

## Cloud Computing PaaS Middleware

Christopher A. Wood  
`www.christopher-wood.com`  
`woodc1@uci.edu`

May 5, 2014

# Agenda

Motivation

Google App Engine Middleware

Front-End

Back-End

- Data Storage

- Process and Application Communication

- Process and Task Management

- Misc Content

Wrapping Up

# Introduction

A PaaS is, by definition, a set of middleware components between applications and the infrastructure on which they run to abstract away:

- ▶ Reliability issues
- ▶ Data redundancy guarantees
- ▶ Load balancing
- ▶ ...


## A New Application Infrastructure

The Google App Engine architecture is designed to provide application:

- ▶ Scalability: load balancing, asynchronous event handling, etc.
- ▶ Reliability: transparent fault tolerance
- ▶ Cost efficiency: application cost dynamic varies based on resource usage

How are traditional web applications built to achieve these three properties?...

## Traditional Enterprise Application/Backend Designs




# GAE Goal

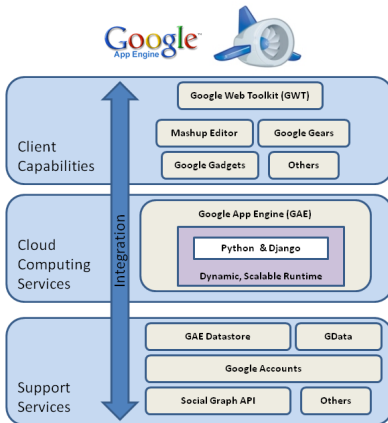
**Question:** How can we get the same *performance* with a significantly *less coupled* backend?

**Answer:** Google App Engine!

# GAE Architecture



## GAE Middleware Role





## “Application-Side” Middleware Abstractions

The application design goals are abstracted by two primary components in the middleware:

- ▶ Front end
  - ▶ Enables “edge caching” at locations close to the user (**load balancing**)
  - ▶ Instances dynamically scale based on the number of incoming requests
- ▶ App server
  - ▶ Runtime environment for application instances (basically, a VM to host your app)

# Backend Middleware

There are many important components of the GAE backend:

- ▶ Data storage, retrieval, and search
- ▶ Communication mechanisms
- ▶ Process management
- ▶ Configuration and management

# Data storage, retrieval, and search

- ▶ Google Cloud SQL
  - ▶ Relational database support for more “legacy” applications
- ▶ Datastore and blobstore vs memcache (and dedicated memcache)
- ▶ Search
  - ▶ Provides a model for indexing (datastore) documents that contain *structured* data, and an API for searching the index
- ▶ ...

# Datastore

A persistent storage for AppEngine

- ▶ BigData/NoSQL data model as opposed to traditional relational data model
- ▶ Queried using GQL (SQL-like query language developed by Google)
- ▶ Significantly better scalability than relational data models
- ▶ Designed as a hierarchy of components: Datastore → MegaStore → BigTable

## Datastore (cont'd)

Handles distribution, replication, and load balancing of data

- ▶ Based on the Paxos consensus algorithm
- ▶ **Scalability**: Automated sharding (i.e., striping) - BigTable
- ▶ **Reliability**: Replication via BigTable and transaction support (with ACID-guarantees) via MegaTable
- ▶ **Performance**: Enhanced lock granularity and co-location of data for better concurrent access (e.g., tables, entities, properties)

## Datastore (cont'd)

- ▶ Provides simple API to read/write entities backed by a powerful query engine and reliable transaction support
- ▶ Stores data as *entities*, which are uniquely identified by keys
  - ▶ Individual data elements are properties of entities (analogous to relation fields)
  - ▶ Entities can have parents and children
  - ▶ Enables a *hierarchical data model*, much like relational data models
  - ▶ Entities without parents are called *root entities*


# Paxos Consensus Algorithm

- ▶ Paxos is a consensus algorithm executed by a set of processes (replicas) to agree on a single value in the presence of failures.
- ▶ Composed of three stages:
  1. Elect a replica to be the coordinator (**we've seen this - distributed elections**).
  2. The coordinator selects a value and broadcasts it to all replicas in a message called the accept message. Other replicas either acknowledge this message or reject it (**we've seen this too - load balancing**).
  3. Once a majority of the replicas acknowledge the coordinator, consensus has been reached, and the coordinator broadcasts a commit message to notify replicas (**yup, we've seen this as well - load balancing again**).
- ▶ Google's implementation transformed a single page of pseudocode into several thousands of lines of C++ code.

# Memcache

An in-memory, key-value data store

- Any serializable object/value can be inserted as a key or value





## Memcache (cont'd)

- ▶ Used for caching (obviously) datastore query results, user session information, etc.
- ▶ Also used for sharing data cross app instances
- ▶ APIs support general (atomic) read/write operations, as well as batch operations (e.g., `writeAll()`, `readAll()`, ...)
- ▶ Memcache **gets** be up to **10x** faster than the datastore **queries**
  - ▶ High memcache hit rates lead to *massive* application performance
- ▶ Dedicated memcaches can be created (for a price) to store data specifically for your application
  - ▶ → no contention among other applications for cache space!

## Memcache (cont'd)

Memcache sample usage:

```
def get_data():  
    data = memcache.get('key')  
    if data is not None:  
        return data  
    else:  
        data = self.query_for_data()  
        memcache.add('key', data, 60)  
        return data
```

# Communication mechanisms


- ▶ Channels
- ▶ Google Cloud Endpoints
- ▶ Mail communication, URL fetch, XMPP (common instant messaging protocol)
  - ▶ Exposed via natural APIs for more “traditional” application or host-to-host (instance-to-instance) communication
  - ▶ Uses Google infrastructure for improved communication efficiency (e.g., for HTTP(s) GET requests issued to URLs)

# Channel Communication

- ▶ Creates a persistent connection between applications and Google servers
- ▶ Applications send messages to JavaScript clients in real time without “polling”
  - ▶ *Enables real-time updates to be pushed to clients*
- ▶ Used when asynchronous events/information needs to be pushed to clients
- ▶ Clients open channel connections to AppServer instances, and servers *push* data to all open channels when needed

## Cloud Endpoint Communication

- ▶ Enables automatic generation of APIs via a set of tools, libraries, and components
  - ▶ API code is *annotated* so that the GAE tools/libraries can automatically generate corresponding APIs for all platforms
- ▶ Used to create a *shared* web backend for web clients and mobile clients (e.g., Android or iOS)
  - ▶ Greatly reduces the amount of platform-specific work that needs to be done



# Process management

Batch processing is handled by two primary components:

- ▶ Task queue
- ▶ Task scheduling

# Task Queue

- ▶ Used to handle user tasks *outside* of a user request (e.g., execute the task asynchronously in the background)
- ▶ Tasks are configured according to a **Task** interface (e.g., Java task instances must implement the Task interface)
- ▶ Queues are configured on a per-application basis using YAML
  - ▶ Both *push* and *pull* queues can be defined
- ▶ Tasks can be enqueued as part of a datastore transaction
- ▶ Follows a “leaky bucket” approach to handle task bandwidth and burstiness submissions to AppServer instances
- ▶ RPC stubs (akin to Java Future objects) are returned when tasks are submitted so that task statuses can be queried

## Task (Job) Scheduler

- ▶ Cron scheduler support for a variety of application languages
- ▶ Jobs can be scheduled to execute at periodic times (e.g., nightly batch processing tasks)
- ▶ GAE provides a tiered price plan for the number of allowable jobs that can be scheduled
- ▶ Jobs are configured with the `cron.xml` file
- ▶ Authorization rules can be built into the cron job



## Configuration and Management

And of course, the GAE middleware provides the following additional features

- ▶ Application identity
- ▶ Remote application access
- ▶ User and application capabilities
- ▶ (custom domain) SSL certificate configuration
- ▶ Traffic splitting (e.g., *roll out* new features over time)
- ▶ User management (e.g., support login using Google accounts or OpenID)

## GAE in One Sentence

GAE extends Google's scalable enterprise backend to your applications with a virtually transparent, flexible, and easy-to-use middleware.

## References

All images taken from Google Developers documentation:

<https://developers.google.com/appengine/features/>